Santhosh Andavar

2818372

CIS 492(Big Data)

04/20/2025

Lab4 part 2

In the second part of the lab, we computed the cosine similarity between 10 selected documents from the State of the Union dataset using the TF-IDF vector space model.

Steps:

- Used TfidfVectorizer from scikit-learn to compute document vectors
- Computed pairwise cosine similarity across documents
- Saved a 10x10 similarity matrix to CSV format

Query: `freedom peace economy`

Top 5 Similar Documents:

- doc5
- doc2
- doc6
- doc3
- doc7

Cosine similarity matrix showing pairwise document similarity across the 10 State of the Union addresses using TF-IDF vectors.

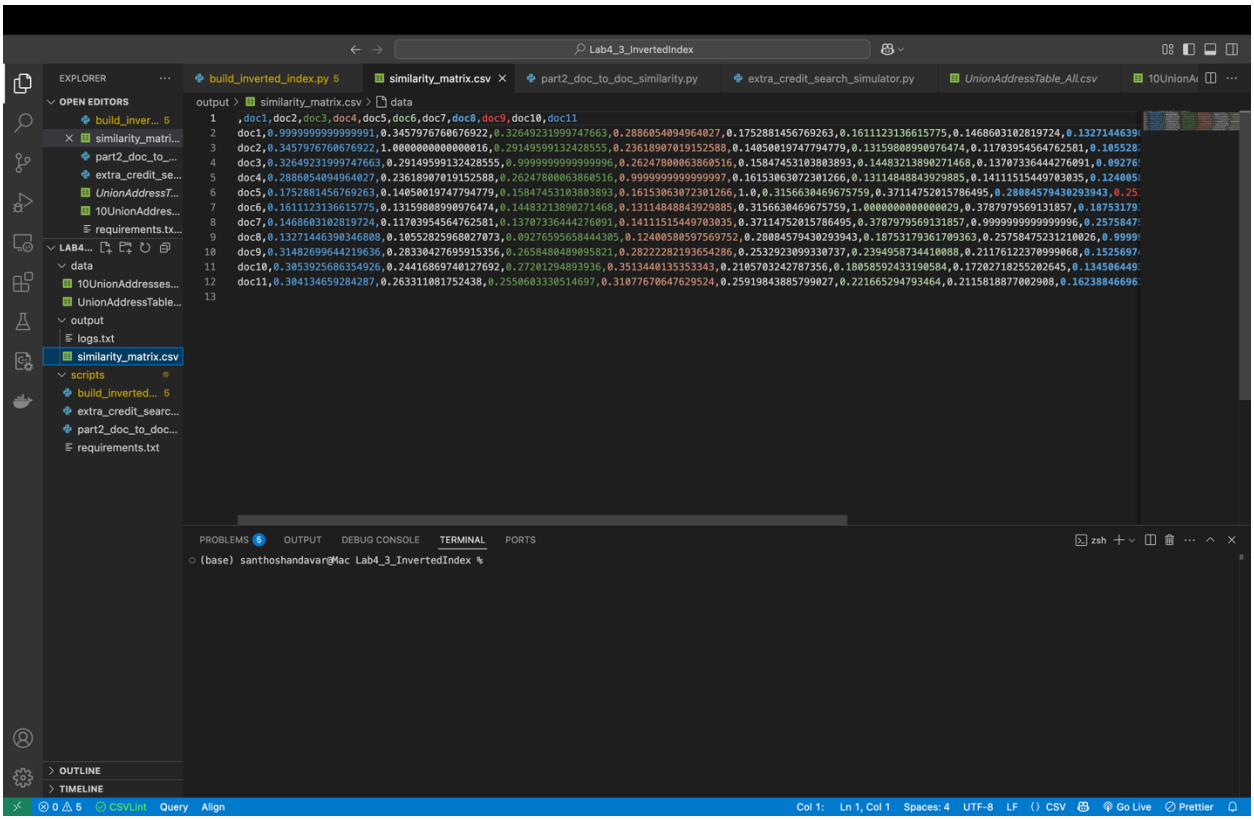

Extra Credit: Web-Based Search Engine with User Input

A Streamlit web application was developed to simulate a real-time search engine. Users can enter a query, and the system returns the top 5 most relevant documents using TF-IDF and cosine similarity.

- Input box captures keywords like "freedom peace economy"
- Matches query vector to document vectors
- Displays most similar speeches and scores

Queries Used for Screenshots:

1. `freedom rights democracy`
2. `economy inflation growth`
3. `war peace treaty`
4. `education science innovation`
5. `tax revenue budget`

Real-time Streamlit app matching user queries against speech documents using cosine similarity.

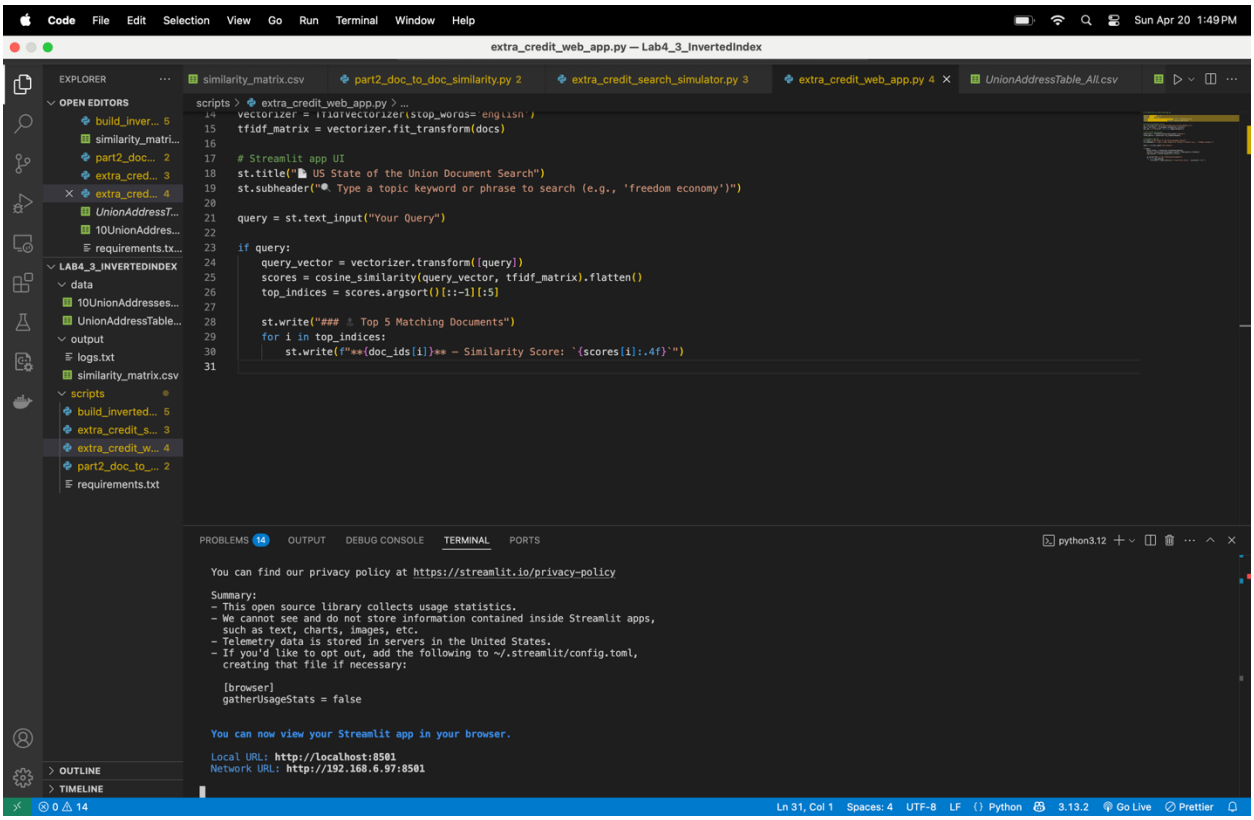

"Freedom rights democracy" — shows documents discussing civil liberties and governance.

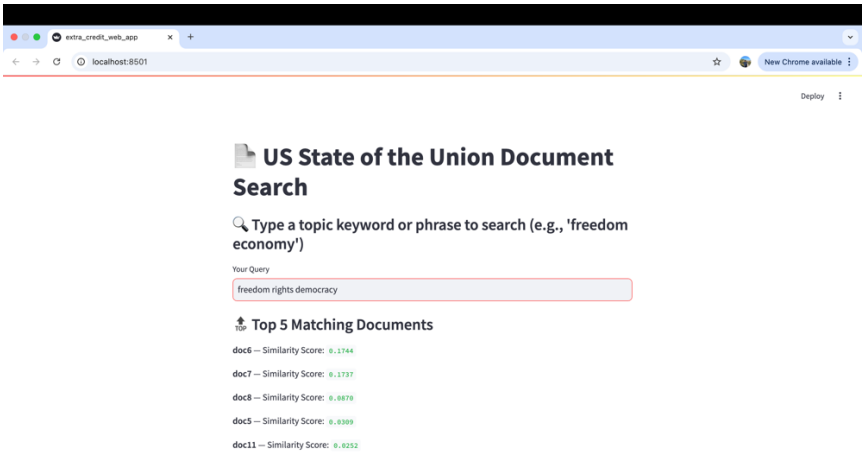

"Economy inflation growth" — returns economic-focused speeches.

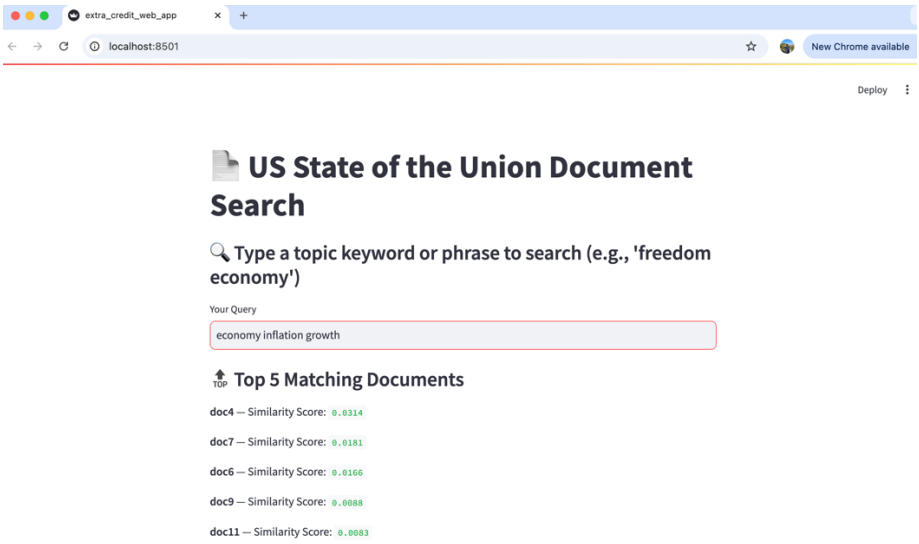

″War peace treaty" — matches documents addressing defense and foreign relations.

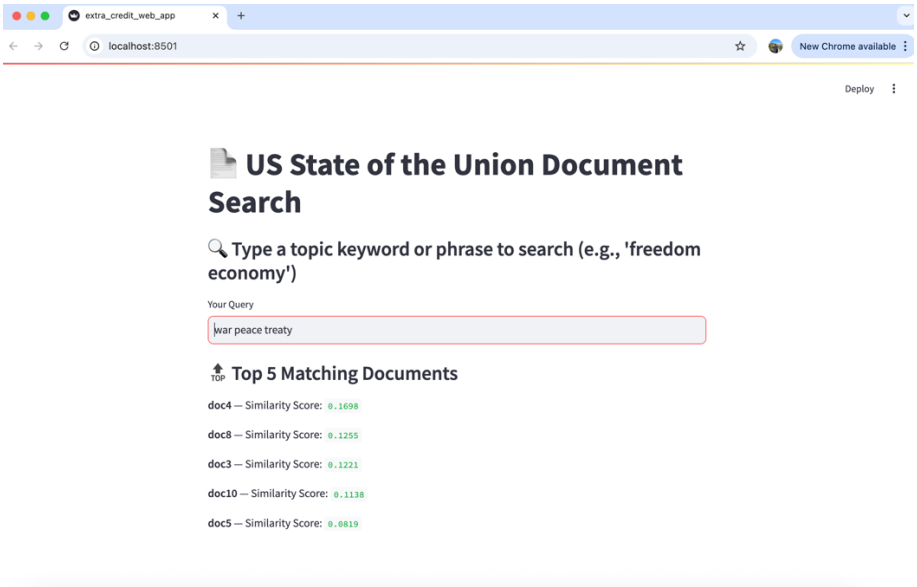

"Education science innovation" — highlights educational and innovation content

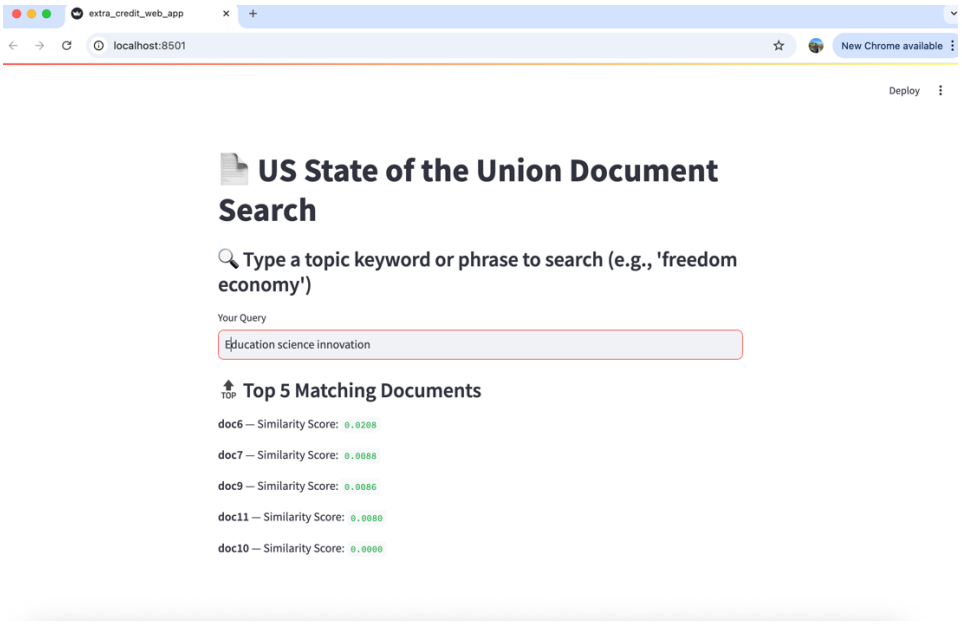

"Tax revenue budget" — captures fiscal discussions in presidential addresses.

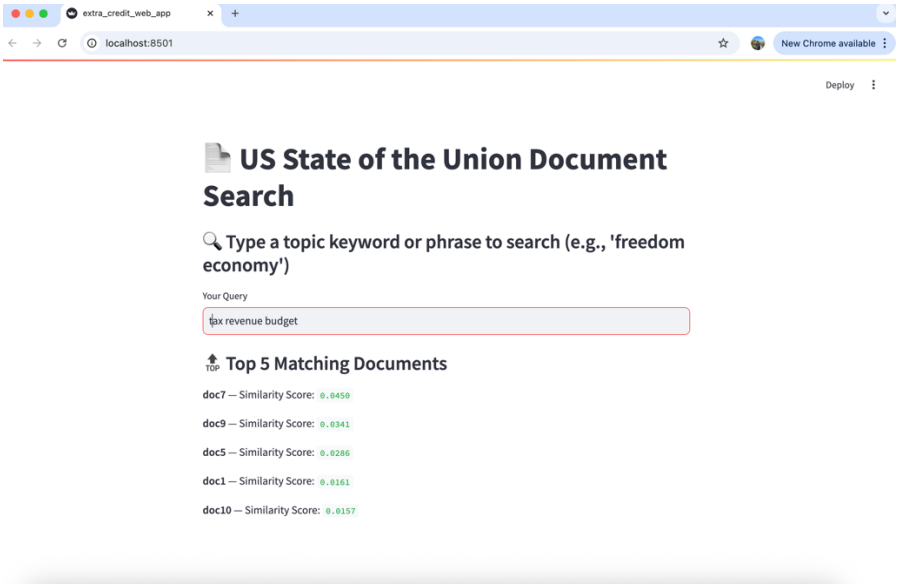

Conclusion:

This lab successfully demonstrated building an inverted index and computing document similarity using TF-IDF and cosine similarity. The integration of NLP techniques and a web-based query system allowed us to explore real-time content-based search with extra credit features, making this a comprehensive document search pipeline.